


## Chapter 2 Application Layer

All material copyright 1996-2016  
© J.F Kurose and K.W. Ross, All Rights Reserved



### *Computer Networking: A Top Down Approach*

7<sup>th</sup> edition

Jim Kurose, Keith Ross  
Pearson/Addison Wesley  
April 2016

Application Layer 2-1

## Chapter 2: outline

### 2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and  
content distribution  
networks

2.7 socket programming  
with UDP and TCP

Application Layer 2-2

## Chapter 2: application layer

### our goals:

- conceptual, implementation aspects of network application protocols
  - transport-layer service models
  - client-server paradigm
  - peer-to-peer paradigm
  - content distribution networks
- learn about protocols by examining popular application-level protocols
  - HTTP
  - FTP
  - SMTP / POP3 / IMAP
  - DNS
- creating network applications
  - socket API

Application Layer 2-3

## Some network apps

- e-mail
- web
- text messaging
- remote login
- P2P file sharing
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- voice over IP (e.g., Skype)
- real-time video conferencing
- social networking
- search
- ...
- ...

Application Layer 2-4

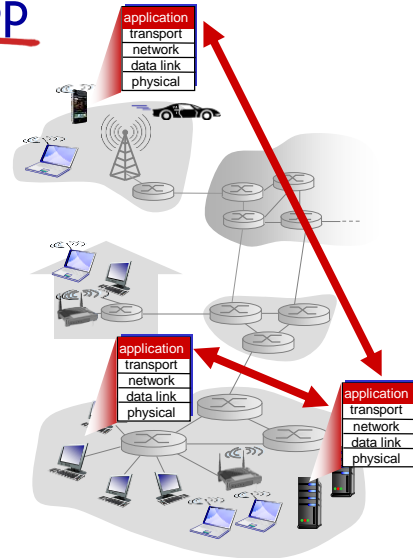
## Creating a network app

### write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

### no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



Application Layer 2-5


## Application architectures

### possible structure of applications:

- client-server
- peer-to-peer (P2P)

Application Layer 2-6

## Client-server architecture



### server:

- always-on host
- permanent IP address
- data centers for scaling


### clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

Application Layer 2-7

## P2P architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management



Application Layer 2-8

# Processes communicating

*process*: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**

clients, servers

*client process*: process that initiates communication


*server process*: process that waits to be contacted

- aside: applications with P2P architectures have client processes & server processes

Application Layer 2-9

# Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
  - sending process shoves message out door
  - sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Application Layer 2-10

## Addressing processes

- to receive messages, process must have *identifier*
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
  - A: no, many processes can be running on same host
- *identifier* includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
  - HTTP server: 80
  - mail server: 25
- to send HTTP message to gaia.cs.umass.edu web server:
  - **IP address**: 128.119.245.12
  - **port number**: 80
- more shortly...

Application Layer 2-11

## App-layer protocol defines

- **types of messages exchanged**,
  - e.g., request, response
- **message syntax**:
  - what fields in messages & how fields are delineated
- **message semantics**
  - meaning of information in fields
- **rules** for when and how processes send & respond to messages
- **open protocols**:
  - defined in RFCs
  - allows for interoperability
  - e.g., HTTP, SMTP
- **proprietary protocols**:
  - e.g., Skype

Application Layer 2-12

## What transport service does an app need?

### data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

### timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

### throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be “effective”
- other apps (“elastic apps”) make use of whatever throughput they get

### security

- encryption, data integrity, ...

Application Layer 2-13

## Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	
interactive games	loss-tolerant	few kbps up	yes, few secs
text messaging	no loss	elastic	yes, 100' s msec yes and no

Application Layer 2-14

## Internet transport protocols services

### TCP service:

- *reliable transport* between sending and receiving process
- *flow control*: sender won't overwhelm receiver
- *congestion control*: throttle sender when network overloaded
- *does not provide*: timing, minimum throughput guarantee, security
- *connection-oriented*: setup required between client and server processes

### UDP service:

- *unreliable data transfer* between sending and receiving process
- *does not provide*: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup,

**Q:** why bother? Why is there a UDP?

Application Layer 2-15

## Internet apps: application, transport protocols

application	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

Application Layer 2-16



## Securing TCP

### TCP & UDP

- no encryption
- cleartext passwords sent into socket traverse Internet in cleartext

### SSL

- provides encrypted TCP connection
- data integrity
- end-point authentication

### SSL is at app layer

- apps use SSL libraries, that “talk” to TCP

### SSL socket API

- cleartext passwords sent into socket traverse Internet encrypted
- see Chapter 8

Application Layer 2-17

## Chapter 2: outline

2.1 principles of network applications

### 2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Application Layer 2-18

# Web and HTTP

*First, a review...*

- **web page** consists of **objects**
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of **base HTML-file** which includes **several referenced objects**
- each object is addressable by a **URL**, e.g.,

`www.someschool.edu/someDept/pic.gif`

host name

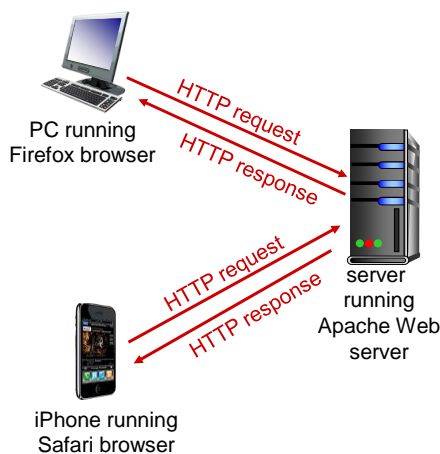
path name

Application Layer 2-19

## HTTP overview

**HTTP: hypertext transfer protocol**

- Web's application layer protocol
- client/server model
  - **client**: browser that requests, receives, (using HTTP protocol) and "displays" Web objects
  - **server**: Web server sends (using HTTP protocol) objects in response to requests



Application Layer 2-20

## HTTP overview (continued)

### *uses TCP:*

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

### *HTTP is “stateless”*

- server maintains no information about past client requests

### *aside*

#### protocols that maintain “state” are complex!

- past history (state) must be maintained
- if server/client crashes, their views of “state” may be inconsistent, must be reconciled

Application Layer 2-21

## HTTP connections

### *non-persistent HTTP*

- at most one object sent over TCP connection
  - connection then closed
- downloading multiple objects required multiple connections

### *persistent HTTP*

- multiple objects can be sent over single TCP connection between client, server

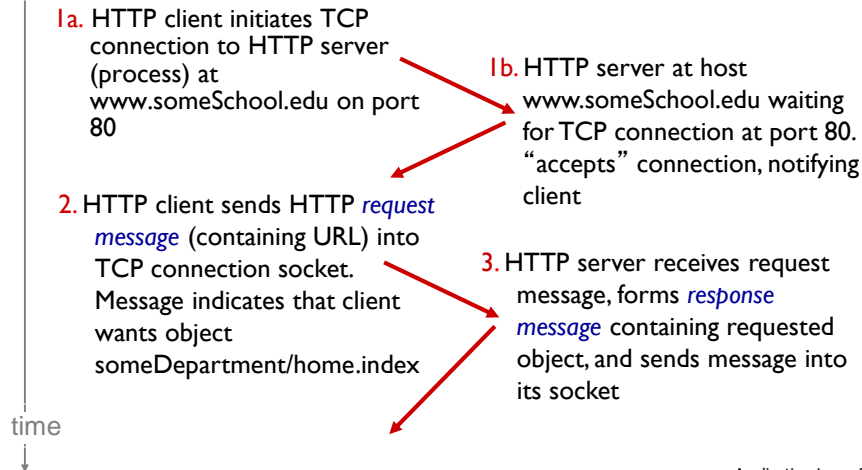
Application Layer 2-22

## Non-persistent HTTP

suppose user enters URL:

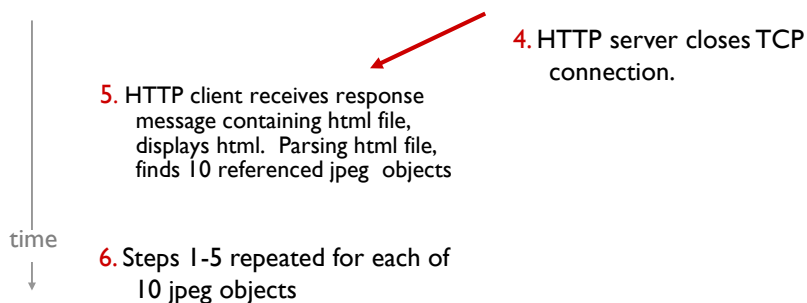
`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)



Application Layer 2-23

## Non-persistent HTTP (cont.)




Application Layer 2-24

## Non-persistent HTTP: response time

**RTT (definition):** time for a small packet to travel from client to server and back

**HTTP response time:**

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =  $2\text{RTT} + \text{file transmission time}$



Application Layer 2-25

## Persistent HTTP

**non-persistent HTTP issues:**

- requires 2 RTTs per object
- OS overhead for each TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

***persistent HTTP:***

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects

Application Layer 2-26

# HTTP request message

- two types of HTTP messages: *request, response*
- HTTP request message:**
  - ASCII (human-readable format)

request line  
(GET, POST,  
HEAD commands)

header  
lines

carriage return,  
line feed at start  
of line indicates  
end of header lines

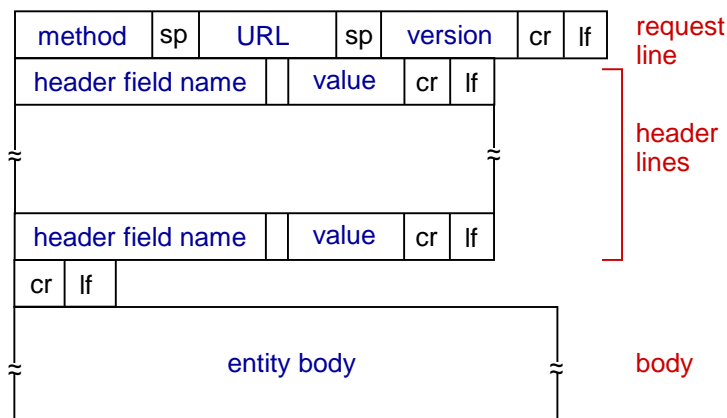
```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Firefox/3.6.10\r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n
Keep-Alive: 115\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character  
line-feed character

\* Check out the online interactive exercises for more  
examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

Application Layer 2-27

# HTTP request message: general format



Application Layer 2-28

## Uploading form input

### POST method:

- web page often includes form input
- input is uploaded to server in entity body

### URL method:

- uses GET method
- input is uploaded in URL field of request line:

`www.somesite.com/animalsearch?monkeys&banana`

Application Layer 2-29

## Method types

### HTTP/1.0:

- GET
- POST
- HEAD
  - asks server to leave requested object out of response

### HTTP/1.1:

- GET, POST, HEAD
- PUT
  - uploads file in entity body to path specified in URL field
- DELETE
  - deletes file specified in the URL field

Application Layer 2-30

## HTTP response message

status line  
(protocol  
status code  
status phrase)

header  
lines

data, e.g.,  
requested  
HTML file

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02
GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-
1\r\n
\r\n
data data data data data ...
```

\* Check out the online interactive exercises for more  
examples: [http://gaia.cs.umass.edu/kurose\\_ross/interactive/](http://gaia.cs.umass.edu/kurose_ross/interactive/)

Application Layer 2-31

## HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:
  - 200 OK**
    - request succeeded, requested object later in this msg
  - 301 Moved Permanently**
    - requested object moved, new location specified later in this msg (Location:)
  - 400 Bad Request**
    - request msg not understood by server
  - 404 Not Found**
    - requested document not found on this server
  - 505 HTTP Version Not Supported**

Application Layer 2-32



## Trying out HTTP (client side) for yourself

1. Telnet to your favorite Web server:

`telnet gaia.cs.umass.edu 80` { opens TCP connection to port 80  
(default HTTP server port)  
at gaia.cs.umass.edu.  
anything typed in will be sent  
to port 80 at gaia.cs.umass.edu

2. type in a GET HTTP request:

`GET /kurose_ross/interactive/index.php HTTP/1.1`  
`Host: gaia.cs.umass.edu` { by typing this in (hit carriage  
return twice), you send  
this minimal (but complete)  
GET request to HTTP server

3. look at response message sent by HTTP server!  
(or use Wireshark to look at captured HTTP request/response)

Application Layer 2-33

## User-server state: cookies

many Web sites use cookies

*four components:*

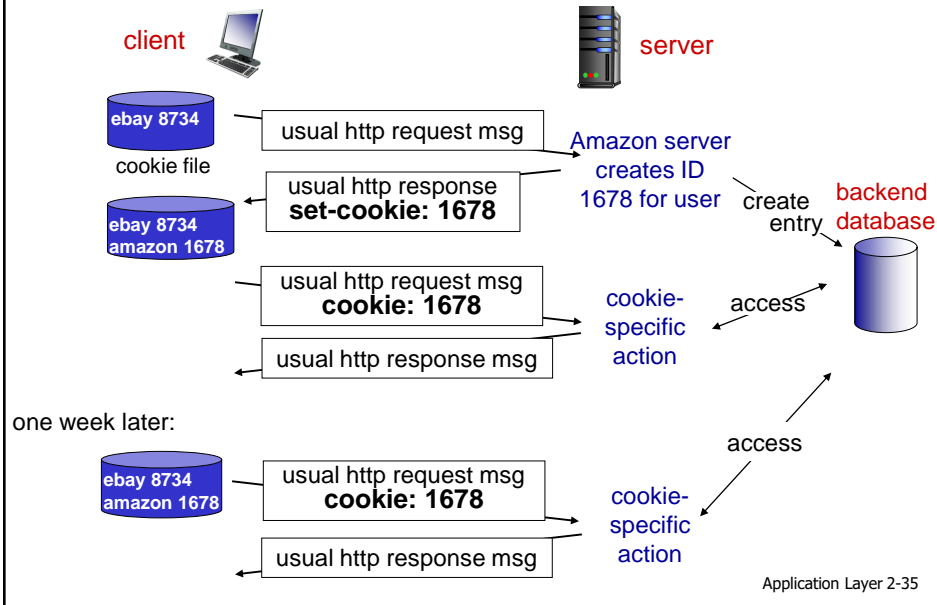
- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

**example:**

- Susan always access Internet from PC
- visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

Application Layer 2-34

## Cookies: keeping “state” (cont.)



## Cookies (continued)

*what cookies can be used for:*

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

*how to keep “state”:*

- protocol endpoints: maintain state at sender/receiver over multiple transactions
- cookies: http messages carry state

aside

*cookies and privacy:*

- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

Application Layer 2-36

## Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version

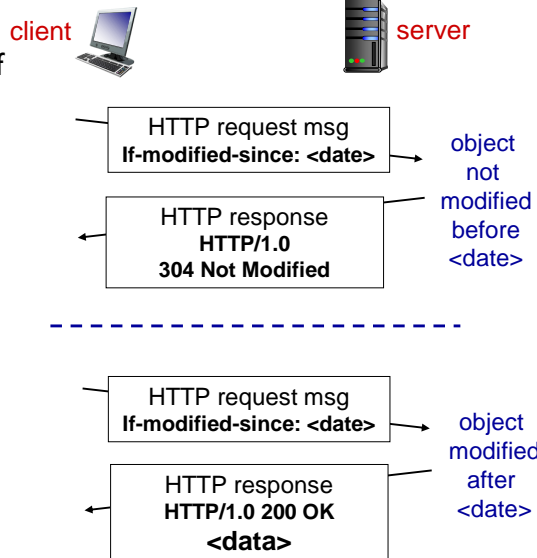
- no object transmission delay
- lower link utilization

- **cache:** specify date of cached copy in HTTP request

If-modified-since: <date>

- **server:** response contains no object if cached copy is up-to-date:

HTTP/1.0 304 Not Modified



Application Layer 2-43

## Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

2.7 socket programming with UDP and TCP

Application Layer 2-56

## DNS: domain name system

*people*: many identifiers:

- SSN, name, passport #

*Internet hosts, routers*:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., www.yahoo.com - used by humans

Q: how to map between IP address and name, and vice versa ?

*Domain Name System*:

- *distributed database* implemented in hierarchy of many *name servers*
- *application-layer protocol*: hosts, name servers communicate to *resolve* names (address/name translation)
  - note: core Internet function, implemented as application-layer protocol
  - complexity at network's “edge”

Application Layer 2-57

## DNS: services, structure

*DNS services*

- hostname to IP address translation
- host aliasing
  - canonical, alias names
- mail server aliasing
- load distribution
  - replicated Web servers: many IP addresses correspond to one name


*why not centralize DNS?*

- single point of failure
- traffic volume
- distant centralized database
- maintenance

*A: doesn't scale!*

Application Layer 2-58

## DNS: a distributed, hierarchical database



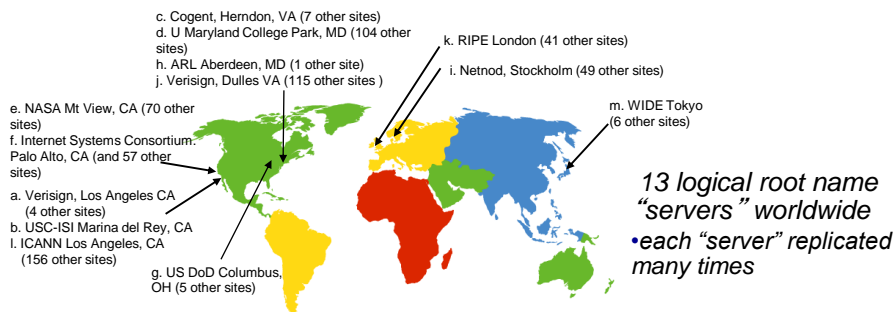
*client wants IP for www.amazon.com; 1<sup>st</sup> approximation:*

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

Application Layer 2-59

## DNS: root name servers

- contacted by local name server that can not resolve name
- root name server:
  - contacts authoritative name server if name mapping not known
  - gets mapping
  - returns mapping to local name server



Application Layer 2-60

## TLD, authoritative servers

### *top-level domain (TLD) servers:*

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

### *authoritative DNS servers:*

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

Application Layer 2-61

## Local DNS name server

- does not strictly belong to hierarchy
- each ISP (residential ISP, company, university) has one
  - also called “default name server”
- when host makes DNS query, query is sent to its local DNS server
  - has local cache of recent name-to-address translation pairs (but may be out of date!)
  - acts as proxy, forwards query into hierarchy


Application Layer 2-62

## DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

### *iterated query:*

- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"




Application Layer 2-63

## DNS name resolution example

### *recursive query:*

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



Application Layer 2-64

## DNS: caching, updating records

- once (any) name server learns mapping, it *caches* mapping
  - cache entries timeout (disappear) after some time (TTL)
  - TLD servers typically cached in local name servers
    - thus root name servers not often visited
- cached entries may be *out-of-date* (best effort name-to-address translation!)
  - if name host changes IP address, may not be known Internet-wide until all TTLs expire
- update/notify mechanisms proposed IETF standard
  - RFC 2136

Application Layer 2-65

## DNS records

*DNS*: distributed database storing resource records (RR)

RR format: (name, value, type, ttl)

### type=A

- **name** is hostname
- **value** is IP address

### type=NS

- **name** is domain (e.g., foo.com)
- **value** is hostname of authoritative name server for this domain

### type=CNAME

- **name** is alias name for some “canonical” (the real) name
- **www.ibm.com** is really **servereast.backup2.ibm.com**
- **value** is canonical name

### type=MX

- **value** is name of mailserver associated with **name**

Application Layer 2-66



## Inserting records into DNS

- example: new startup “Network Utopia”
- register name networkutopia.com at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative name server (primary and secondary)
  - registrar inserts two RRs into .com TLD server:  
(networkutopia.com, dns1.networkutopia.com, NS)  
(dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server type A record for www.networkutopia.com; type MX record for networkutopia.com

Application Layer 2-69

## Attacking DNS

### DDoS attacks

- bombard root servers with traffic
  - not successful to date
  - traffic filtering
  - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
  - potentially more dangerous

### redirect attacks

- man-in-middle
  - Intercept queries
- DNS poisoning
  - Send bogus replies to DNS server, which caches

### exploit DNS for DDoS

- send queries with spoofed source address: target IP
- requires amplification

Application Layer 2-70

## Chapter 2: outline

2.1 principles of network applications

2.2 Web and HTTP

2.3 electronic mail

- SMTP, POP3, IMAP

2.4 DNS

2.5 P2P applications

2.6 video streaming and content distribution networks

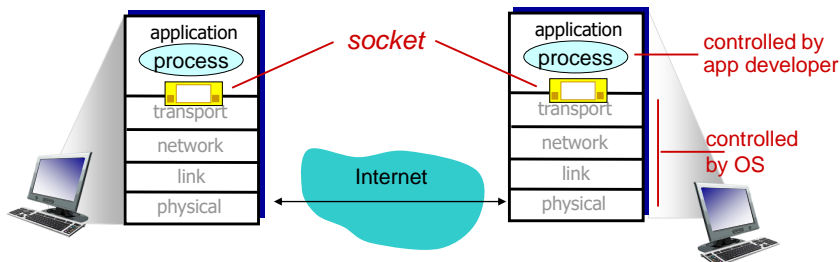
2.7 socket programming with UDP and TCP

Application Layer 2-94

## Socket programming

**goal:** learn how to build client/server applications that communicate using sockets

**socket:** door between application process and end-end-transport protocol



Application Layer 2-95

## Socket programming

*Two socket types for two transport services:*

- **UDP:** unreliable datagram
- **TCP:** reliable, byte stream-oriented

*Application Example:*

1. client reads a line of characters (data) from its keyboard and sends data to server
2. server receives the data and converts characters to uppercase
3. server sends modified data to client
4. client receives modified data and displays line on its screen

Application Layer 2-96

## Socket programming **with UDP**

**UDP: no “connection” between client & server**

- no handshaking before sending data
- sender explicitly attaches IP destination address and port # to each packet
- receiver extracts sender IP address and port# from received packet

**UDP: transmitted data may be lost or received out-of-order**

**Application viewpoint:**

- UDP provides *unreliable* transfer of groups of bytes (“datagrams”) between client and server

Application Layer 2-97

## Client/server socket interaction: UDP

### server (running on serverIP)

create socket, port= x:  
`serverSocket =  
 socket(AF_INET,SOCK_DGRAM)`

read datagram from  
`serverSocket`

write reply to  
`serverSocket`  
 specifying  
 client address,  
 port number

### client

create socket:  
`clientSocket =  
 socket(AF_INET,SOCK_DGRAM)`

Create datagram with server IP and  
 port=x; send datagram via  
`clientSocket`

read datagram from  
`clientSocket`

close  
`clientSocket`

Application 2-98

## Example app: UDP client

### *Python UDPClient*

include Python's socket  
 library

from socket import \*  
 serverName = 'hostname'  
 serverPort = 12000

create UDP socket for  
 server

`clientSocket = socket(AF_INET,  
 SOCK_DGRAM)`

get user keyboard  
 input

`message = input('Input lowercase sentence:')`

Attach server name, port to  
 message; send into socket

`clientSocket.sendto(message.encode(),  
 (serverName, serverPort))`

read reply characters from  
 socket into string

`modifiedMessage, serverAddress =  
 clientSocket.recvfrom(2048)`

print out received string  
 and close socket

`Print(modifiedMessage.decode())  
 clientSocket.close()`

Application Layer 2-99

## Example app: UDP server

### *Python UDPServer*

```
from socket import *
serverPort = 12000

create UDP socket → serverSocket = socket(AF_INET, SOCK_DGRAM)
bind socket to local port → serverSocket.bind(("", serverPort))
number 12000
print ("The server is ready to receive")

loop forever → while True:
    Read from UDP socket into → message, clientAddress = serverSocket.recvfrom(2048)
    message, getting client's → modifiedMessage = message.decode().upper()
    address (client IP and port)
    send upper case string → serverSocket.sendto(modifiedMessage.encode(),
    back to this client → clientAddress)
```

Application Layer 2-100

## Socket programming *with TCP*

### *client must contact server*

- server process must first be running
- server must have created socket (door) that welcomes client's contact

### *client contacts server by:*

- Creating TCP socket, specifying IP address, port number of server process
- *when client creates socket:* client TCP establishes connection to server TCP

- when contacted by client, *server TCP creates new socket* for server process to communicate with that particular client
  - allows server to talk with multiple clients
  - source port numbers used to distinguish clients (more in Chap 3)

### *application viewpoint:*

TCP provides reliable, in-order byte-stream transfer ("pipe") between client and server

Application Layer 2-101

## Client/server socket interaction: TCP

### server (running on `hostid`)

### client

create socket,  
port=`x`, for incoming  
request:  
`serverSocket = socket()`

wait for incoming  
connection request  
`connectionSocket =`  
`serverSocket.accept()`

read request from  
`connectionSocket`

write reply to  
`connectionSocket`

close  
`connectionSocket`

create socket,  
connect to `hostid`, port=`x`  
`clientSocket = socket()`

send request using  
`clientSocket`

read reply from  
`clientSocket`

close  
`clientSocket`

**TCP**  
connection setup

Application Layer 2-102

## Example app: TCP client

### *Python TCPClient*

```
from socket import *
```

```
serverName = 'servername'
```

```
serverPort = 12000
```

create TCP socket for  
server, remote port 12000

```
clientSocket = socket(AF_INET, SOCK_STREAM)
```

```
clientSocket.connect((serverName, serverPort))
```

```
sentence = input('Input lowercase sentence:')
```

No need to attach server  
name, port

```
clientSocket.send(sentence.encode())
```

```
modifiedSentence = clientSocket.recv(1024)
```

```
print('From Server:', modifiedSentence.decode())
```

```
clientSocket.close()
```

Application Layer 2-103

## Example app: TCP server

### *Python TCPServer*

```
from socket import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_STREAM)
serverSocket.bind(('', serverPort))
serverSocket.listen(1)
print('The server is ready to receive')
while True:
    connectionSocket, addr = serverSocket.accept()
    sentence = connectionSocket.recv(1024).decode()
    capitalizedSentence = sentence.upper()
    connectionSocket.send(capitalizedSentence.
                           encode())
    connectionSocket.close()
```

create TCP welcoming socket →

server begins listening for incoming TCP requests →

loop forever →

server waits on accept() for incoming requests, new socket created on return →

read bytes from socket (but not address as in UDP) →

close connection to this client (but *not* welcoming socket) →

Application Layer 2-104

## Chapter 2: summary

*our study of network apps now complete!*

- application architectures
  - client-server
  - P2P
- application service requirements:
  - reliability, bandwidth, delay
- Internet transport service model
  - connection-oriented, reliable: TCP
  - unreliable, datagrams: UDP
- specific protocols:
  - HTTP
  - SMTP, POP, IMAP
  - DNS
  - P2P: BitTorrent
- video streaming, CDNs
- socket programming: TCP, UDP sockets

Application Layer 2-105

## Chapter 2: summary

*most importantly: learned about protocols!*

- typical request/reply message exchange:
  - client requests info or service
  - server responds with data, status code
- message formats:
  - *headers*: fields giving info about data
  - *data*: info(payload) being communicated

### *important themes:*

- control vs. messages
  - in-band, out-of-band
- centralized vs. decentralized
- stateless vs. stateful
- reliable vs. unreliable message transfer
- “complexity at network edge”

Application Layer 2-106